

Compulsory exercise 2: Group 13

TMA4315: Generalized linear (mixed) models H2023

Benjamin Sigbjørnsen, Johan Vik Mathisen, and Martinius Theodor Singdahlsen

09 October, 2023

Packages:

```
library(car)
library(ggplot2)
library(GGally)
library(mylm)
library(reshape2)
library(kableExtra)
library(dplyr)

set.seed(42)
```

Part 2)

```
# Reading the data
rm(list = ls())
filepath <- "https://www.math.ntnu.no/emner/TMA4315/2023h/eliteserien2023.csv"
eliteserie <- read.csv(file = filepath)

NGames <- table(c(eliteserie$home[!is.na(eliteserie$yh)], eliteserie$away[!is.na(eliteserie$yh)]))
RangeofGames <- range(NGames)
```

a)

To test the assumption of independence, we have to perform the test

H_0 : Independence between the goals made by the home and away teams H_1 : Dependence between the goals made by the home and away teams

We test using a Pearson's χ^2 test on the contingency table of the goals

```
# Creating contingency table
contingencyTab = table(eliteserie$yh, eliteserie$ya)
# Testing the independence between the goals made by the home and away teams
chisq.test(contingencyTab)

##
## Pearson's Chi-squared test
##
## data:  contingencyTab
## X-squared = 33.527, df = 35, p-value = 0.5393
```

The p-value of the test is 0.54, and we therefore keep our null hypothesis using every reasonable significance level.

b)

The preliminary ranking for the season as of October 1st is presented in table 1.

```
tableFunc = function(eliteserie) {
  # Removing matches that are not played
  eliteserie = eliteserie[!is.na(eliteserie$yh), ]

  # Points home team
  eliteserie$ph = 0
  eliteserie[eliteserie$yh > eliteserie$ya, ]$ph = 3
  eliteserie[eliteserie$yh == eliteserie$ya, ]$ph = 1

  # Points away team
  eliteserie$pa = 1
  eliteserie[eliteserie$ph != eliteserie$pa, ]$pa = 3 - eliteserie[eliteserie$ph !=
    eliteserie$pa, ]$ph

  # Creating table
  table = data.frame(team = c(eliteserie$home[1:8], eliteserie$away[1:8]), points = 0,
    goal_difference = 0)

  # Giving teams points
  for (i in 1:nrow(eliteserie)) {
    teamHome = eliteserie$home[i]
    teamAway = eliteserie$away[i]
    table[table$team == teamHome, ]$points = table[table$team == teamHome, ]$points +
      eliteserie$ph[i]
    table[table$team == teamAway, ]$points = table[table$team == teamAway, ]$points +
      eliteserie$pa[i]
    table[table$team == teamHome, ]$goal_difference = table[table$team == teamHome,
      ]$goal_difference + eliteserie$yh[i] - eliteserie$ya[i]
    table[table$team == teamAway, ]$goal_difference = table[table$team == teamAway,
      ]$goal_difference + eliteserie$ya[i] - eliteserie$yh[i]
  }

  # Sorting by points, and then goal difference
  table = table[order(table$points, table$goal_difference, decreasing = TRUE),
    ]
  return(table)
}
table = tableFunc(eliteserie)
rownames(table) <- 1:nrow(table)
table %>%
  kbl(caption = "Preliminary ranking for the season as of October 1st \\label{tab:prelim_rank}") %>%
  kable_classic(full_width = F, html_font = "Cambria")
```

c)

In this exercise we implement a function estimating the intercept, home advantage and strength parameters for each team in Tippeligaen 2023 based on the 192 games played as of October 1.

Table 1: Preliminary ranking for the season as of October 1st

team	points	goal_difference
Viking	51	21
Bodø/Glimt	49	28
Tromsø	48	11
Brann	42	14
Molde	41	24
Lillestrøm	36	7
Sarpsborg 08	34	5
Odd	30	-3
Rosenborg	29	-6
Strømsgodset	27	-3
HamKam	24	-15
Vålerenga	21	-8
Sandefjord Fotball	21	-9
Haugesund	21	-14
Stabæk	17	-16
Aalesund	12	-36

We build our design matrix X as explained by the hint in the project description.

```
# Removing nans.
eliteserie_clean <- na.omit(eliteserie)
# Extracting nr of games played so far this season.
games_played = dim(eliteserie_clean)[1]
# Extracting team names.
teams = as.array(unique(eliteserie$home))

# Initialize the design matrix X, which is a 2*192 x 18 matrix. Each row in X
# corresponds to a played game. Once from each side. Therefore 2*games_played
# rows.

# Initialize all zeros matrix
X = matrix(0, nrow = games_played * 2, ncol = 18)

# Set proper feature names
colnames(X) = c("Intercept", "HomeAdvantage", teams)

# First we alter the design matrix for all 'home games'
for (i in 1:games_played) {
  # Intercept
  X[i, 1] = 1

  # HomeAdvantage = 1
  X[i, 2] = 1

  # x_homeTeam = 1
  X[i, match(eliteserie_clean[i, 2], colnames(X))] = 1

  # x_awayTeam = -1
  X[i, match(eliteserie_clean[i, 3], colnames(X))] = -1
}
```

```

}

# Then alter design matrix for all 'away games'
for (i in 1:games_played) {

  # Intercept
  X[games_played + i, 1] = 1

  # Same as above, but inverted:

  # x_home = 0 from initialization.

  # x_homeTeam = -1
  X[games_played + i, match(eliteserie_clean[i, 2], colnames(X))] = -1

  # x_awayTeam is set to 1
  X[games_played + i, match(eliteserie_clean[i, 3], colnames(X))] = 1
}

# Response vector, all goals.
Y = c(eliteserie_clean$yh, eliteserie_clean$ya)
Y = matrix(Y, length(Y), 1)

# Removed to make mle solution unique and optim to work.
X <- X[, -which(colnames(X) == "Bodø/Glimt")]

```

For parameter estimation we use maximum likelihood estimation. For a Poisson regression model we know that the MLE does not have a closed form solution, hence we turn to numerics in order to obtain our estimates.

Since our 16 team covariates sum to 0, the MLE has infinitely many solutions. We do as proposed by the exercise text and force $x_{\text{Bodø/Glimt}}$ to be zero, leaving us with a linearly independent set of covariates. In our model the strength parameter for Bodø/Glimt will then be the reference value 0.

The log likelihood of a Poisson sample is

$$l(\beta) = \sum_{i=1}^n [y_i \ln(\lambda_i) - \lambda_i - \ln(y_i!)] .$$

As we use the R function `Optim`, and this by default minimizes, we implement the negative log likelihood below.

By including the data-dependent term $\ln(y_i!)$ in our calculations we got extremely accurate results, and therefore choose not to include gradient information to further increase accuracy.

```

neg_log_likelihood_poisson <- function(beta) {
  eta <- X %*% beta
  lambda <- exp(eta)
  return(-(t(Y) %*% eta - sum(lambda - log(factorial(Y)))))
}

my_poisson <- function(X, Y) {
  "
X: Design matrix
Y: Response (poisson distributed)

```

Table 2: Power ranking based on Poisson regresson coefficients.

Names	Estimate
Bodø/Glimt	0.00000
Molde	-0.10314
Viking	-0.12260
Brann	-0.24826
Tromsø	-0.26013
Lillestrøm	-0.32865
Sarpsborg 08	-0.34396
Odd	-0.44285
Strømsgodset	-0.47908
Sandefjord Fotball	-0.50867
Rosenborg	-0.51028
Vålerenga	-0.52032
Haugesund	-0.57215
Stabæk	-0.63302
HamKam	-0.63641
Aalesund	-0.91919

```

"
  # Initial value for beta.
  beta_init <- rep(0.5, 17)

  beta <- optim(beta_init, neg_log_likelihood_poisson, method = "BFGS")$par
  result <- data.frame(Names = colnames(X), Estimate = round(beta, 5))
  return(result)
}
result <- my_poisson(X, Y)

# Remove intercept and homeAdvantage.
power_ranking <- result[-(1:2), ]
# Adding Bodø/Glimt back
power_ranking <- rbind(power_ranking, list(Names = "Bodø/Glimt", Estimate = 0))

# Order teams by power parameter.
power_ranking <- power_ranking[order(power_ranking[, 2], decreasing = TRUE), ]
rownames(power_ranking) <- 1:nrow(power_ranking)

# Power rankings
power_ranking %>%
  kbl(caption = "Power ranking based on Poisson regresson coefficients.\\label{tab:power_rank}") %>%
  kable_classic(full_width = F, html_font = "Cambria")

# Parameter estimates
result %>%
  kbl(caption = "Estimates for the regresson coefficients.\\label{tab:reg_coeff}") %>%
  kable_classic(full_width = F, html_font = "Cambria")

```

Preliminary ranking vs strength parameter ranking When we compare the ranking based on the estimated strengths presented in table 2 to the preliminary standings in table 1, we observe that the strength

Table 3: Estimates for the regresssion coefficients.

Names	Estimate
Intercept	0.16644
HomeAdvantage	0.35856
Rosenborg	-0.51028
Aalesund	-0.91919
HamKam	-0.63641
Sarpsborg 08	-0.34396
Stabæk	-0.63302
Tromsø	-0.26013
Brann	-0.24826
Lillestrøm	-0.32865
Viking	-0.12260
Haugesund	-0.57215
Molde	-0.10314
Odd	-0.44285
Sandefjord Fotball	-0.50867
Strømsgodset	-0.47908
Vålerenga	-0.52032

parameter follows the goal difference rather than the points. This is not very surprising as our model only takes goal difference into account.

Our results compared to glm() From the summary of the `glm()` model below and our estimated regression coefficients presented in table 3, we see that the estimated coefficients are identical. This leaves us confident in the correctness of our implementation.

```
##
## Call:
## glm(formula = Y ~ -1 + X, family = "poisson")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2652  -1.1581  -0.1047   0.5316   2.4022
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## XIntercept          0.16644    0.06849   2.430 0.015091 *
## XHomeAdvantage       0.35856    0.08745   4.100 4.13e-05 ***
## XRosenborg          -0.51028    0.16443  -3.103 0.001914 **
## XAalesund            -0.91919    0.16662  -5.517 3.46e-08 ***
## XHamKam              -0.63641    0.16580  -3.838 0.000124 ***
## XSarpsborg 08        -0.34396    0.16839  -2.043 0.041089 *
## XStabæk              -0.63302    0.16909  -3.744 0.000181 ***
## XTromsø              -0.26013    0.16452  -1.581 0.113831
## XBrann                -0.24826    0.16821  -1.476 0.139978
## XLillestrøm          -0.32865    0.16986  -1.935 0.053014 .
## XViking               -0.12260    0.16428  -0.746 0.455518
## XHaugesund           -0.57215    0.16436  -3.481 0.000499 ***
## XMolde                -0.10314    0.16687  -0.618 0.536537
## XOdd                  -0.44285    0.16478  -2.688 0.007197 **
## XSandefjord Fotball -0.50867    0.17038  -2.986 0.002830 **
```

```
## XStrømsgodset      -0.47908    0.17080   -2.805 0.005032 **
## XVålerenga        -0.52032    0.16594   -3.136 0.001715 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 555.74  on 358  degrees of freedom
## Residual deviance: 397.07  on 341  degrees of freedom
## AIC: 1102
##
## Number of Fisher Scoring iterations: 5
```

d)

In this exercise we estimate the rankings by means of simulation, based on the estimated strengths found in the previous task. We do this by simulating the remaining 61 matches 1000 times.

In order to obtain the probabilities of the different outcomes of the remaining matches we observe that team A winning over team B is the same as the score Y_A being greater than the score Y_B . From our model-assumptions these variables are independent of each other. From this we obtain that

$$\begin{aligned} P(A \text{ wins over } B) &= P(A \text{ scores more than } B) \\ &= \sum_{k=1}^{\infty} P(Y_A = k, Y_B < k) \\ &= \sum_{k=1}^{\infty} \sum_{j=1}^k P(Y_A = k)P(Y_B = j) \end{aligned}$$

Similarly for draw and loss we have

$$P(A \text{ loses to } B) = \sum_{k=1}^{\infty} \sum_{j=1}^k P(Y_B = k)P(Y_A = j),$$

and

$$P(\text{Draw}) = \sum_{k=0}^{\infty} P(Y_A = k)P(Y_B = k).$$

Let A be the home team. Then from our model we have that Y_A is Poisson with

$$\ln E(Y_A) = \ln(\lambda_A) = \eta_A = \beta_0 + \beta_{\text{home}} + \beta_A - \beta_B$$

and Y_B is Poisson with

$$\ln E(Y_B) = \ln(\lambda_B) = \eta_B = \beta_0 - \beta_A + \beta_B$$

when A plays against B . From this we can calculate the desired probabilities.

The sums involved when calculating the probabilities converge extremely quickly, even for $\lambda_{\text{Bodø/Glimt}} = 4.238418$, which is the largest of all teams, the summands are of magnitude less than 10^{-8} for $k = 20$. Thus we take the liberty to say that the sums have converged when $k = 20$.

```

predict_season <- function(Unplayed, Estimates) {
  "
  Unplayed: home, away
  Estimates: df of poisson regression estimates with names.

  returns: df with match outcome probabilities. Home, Away, Win, Draw, Loss
  "
  Predict <- data.frame(Home = character(0), Away = character(0), Win = numeric(0),
    Draw = numeric(0), Loss = numeric(0))

  for (i in 1:nrow(Unplayed)) {
    home <- Unplayed$home[i]
    away <- Unplayed$away[i]
    p_w <- prob_win(home, away, Estimates)
    p_d <- prob_draw(home, away, Estimates)
    p_l <- prob_loss(home, away, Estimates)
    data <- list(Home = home, Away = away, Win = p_w, Draw = p_d, Loss = p_l)
    Predict <- rbind(Predict, data)
  }
  return(Predict)
}

sim <- function(Predict) {
  "
  Uses Predict df to simulate the remaining season
  "
  outcome <- apply(Predict[, -(1:2)], 1, function(row) {
    return(sample(c("Win", "Draw", "Loss"), 1, prob = row))
  })
  return(outcome)
}

predict_to_table <- function(Predict, outcome, current_standings) {
  "
  Updates the standings according to the
  simulation.

  Predict: Home, Away, probabilities,
  outcome: Sim; result of simulation

  current_standings: Team, points
  "
  Predict <- cbind(Predict, Sim = outcome)

  table <- current_standings
  for (i in 1:nrow(Predict)) {
    if (Predict$Sim[i] == "Win") {
      table[table$team == Predict$Home[i], 2] <- table[table$team == Predict$Home[i],
        2] + 3
    } else if (Predict$Sim[i] == "Loss") {
      table[table$team == Predict$Away[i], 2] <- table[table$team == Predict$Away[i],
        2] + 3
    } else {
      # Draw
    }
  }
}

```



```

        table[table$team == Predict$Home[i], 2] <- table[table$team == Predict$Home[i],
          2] + 1
        table[table$team == Predict$Away[i], 2] <- table[table$team == Predict$Away[i],
          2] + 1
      }
    }
    table <- table[order(table$points, decreasing = TRUE), ]
    return(table)
  }
}

sim_season_finish <- function(number_of_sims = 1000, current_standing, Unplayed,
  Estimate) {
  "
  Combines the above functions to simulate the season number_of_sims times.
  "
  Predict <- predict_season(Unplayed, Estimate)
  outcome <- sim(Predict)
  table <- predict_to_table(Predict, outcome, current_standing)
  for (i in 1:number_of_sims) {
    outcome <- sim(Predict)
    table_1 <- predict_to_table(Predict, outcome, current_standing)
    table <- inner_join(table, table_1, by = "team")
  }
  # Renames the columns such that results from simulation i is in column 'i'
  colnames(table)[2:ncol(table)] <- 1:(ncol(table) - 1)
  as.data.frame(table)
  return(table)
}

get_summary_data <- function(sim_result) {
  "
  Counts the number of times each team finishes at the top of the table.
  Calculates mean points obtained for each team.
  Calculates standard deviation of points for each team.
  "
  N <- ncol(sim_result) - 1 #not including team column. This is the number of sims.
  series_win <- data.frame(`SeasonsWon(%)` = rep(0, length(sim_result$team)))
  rownames(series_win) <- sim_result$team
  for (i in 2:ncol(res)) {
    series_win[res[which.max(sim_result[, i]), 1], ] <- series_win[sim_result[which.max(res[,
      i]), 1], ] + 1
  }
  mean <- apply(sim_result[, -1], 1, function(row) {
    return(mean(row))
  })
  std.d <- apply(sim_result[, -1], 1, function(row) {
    return(sd(row))
  })
  summary_data <- list(Mean = mean, Std.dev = std.d, SeasonsWon = series_win/N)
  summary_data <- as.data.frame(summary_data)

  return(summary_data[order(-summary_data$Mean), ])
}

```

Table 4: Summary data from 1000 end of season simulations.

	Mean	Std.dev	SeasonsWon...
Bodø/Glimt	66.45854	3.236129	0.4685315
Viking	66.28172	3.465913	0.5174825
Tromsø	59.75125	3.124910	0.0129870
Molde	55.46853	3.527784	0.0009990
Brann	54.10989	3.667412	0.0000000
Lillestrøm	47.99800	3.487979	0.0000000
Sarpsborg 08	44.83017	3.240236	0.0000000
Odd	39.75125	3.230025	0.0000000
Rosenborg	36.88312	3.145048	0.0000000
Strømsgodset	36.37063	3.367120	0.0000000
Sandefjord Fotball	32.34466	3.502298	0.0000000
HamKam	31.36364	3.176104	0.0000000
Vålerenga	30.93906	3.550392	0.0000000
Haugesund	30.18482	3.149732	0.0000000
Stabæk	25.09291	3.400347	0.0000000
Aalesund	15.29371	2.385718	0.0000000

In table 4 we see mean number of points and the associated standard deviations, together with the proportion of seasons won by each of the teams in the 1000 simulated season endings. The low standard deviation stands to us out as the least true to reality for a game of this nature. In figure 2 and 3 we see the empirical point distributions for each of the teams, both jointly and separately.

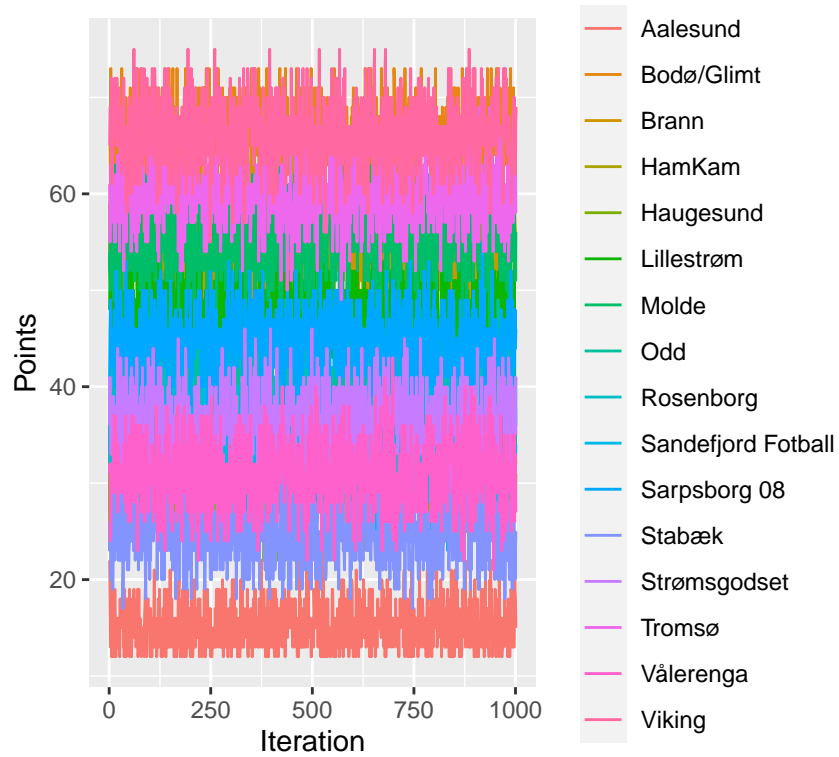


Figure 1: Simulated points at the end of the season, given the standings as of Oct. 1.

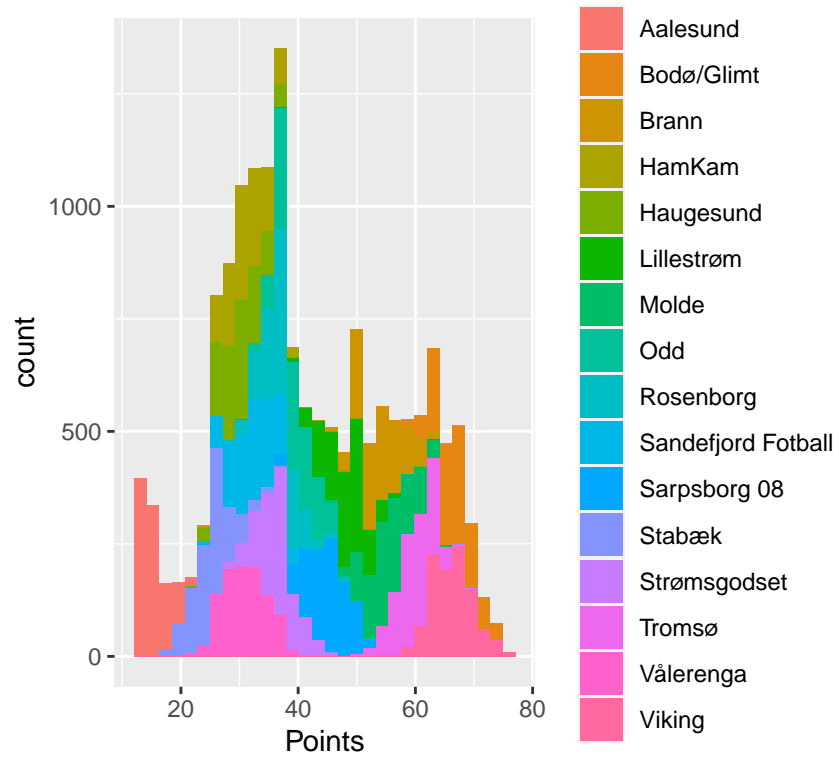


Figure 2: End of season point distribution after 1000 simulations, given the standings as of Oct. 1.




Figure 3: End of season point distribution after 1000 simulations for the individual teams, given the standings as of Oct. 1.